

TURING.JL

<https://github.com/yebai/Turing.jl>

A fresh approach to probabilistic programming in Julia

👤 Kai Xu 📅 JuliaCon, 22 June 2017

👥 Hong Ge & Zoubin Ghahramani

🎓 Machine Learning Group, University of Cambridge

OUTLINE

- What is probabilistic programming?
- Why Julia?
- Key features
- How Turing.jl works?
- Live demo(s)
- Q&A

WHAT IS PROBABILISTIC PROGRAMMING?

- Probabilistic programming is general-purpose programming with intrinsic support of non-deterministic statements.
- Such programming languages are called probabilistic programming languages (PPLs).
- One of the main applications of PPLs is probabilistic modelling, a popular field of machine learning.

- Two types of PPLs
 - Standalone:
 - BUGS
 - Stan
 - Embedded-in:
 - Probabilistic C in C
 - Anglican in Clojure
 - Edward in Python
 - WebPPL in JS
 - Turing.jl in ★Julia★

$$\sigma^2 \sim \text{Inv-Gamma}(2, 3)$$

$$\mu \sim \text{Normal}(0, \sigma)$$

$$x_1, x_2 \sim \text{Normal}(\mu, \sigma)$$

```
@model gdemo(x) = begin
    σ² ~ InverseGamma(2, 3)
    μ ~ Normal(0, sqrt(σ²))
    x[1] ~ Normal(μ, sqrt(σ²))
    x[2] ~ Normal(μ, sqrt(σ²))
    σ², μ
end
```

Code 1: Gaussian model with conjugate priors

Looking closely, we can see the probabilistic features we mentioned before in Turing.jl ...

```
using Turing

@model gdemo(x) = begin
     $\sigma^2$  ~ InverseGamma(2, 3)
     $\mu$  ~ Normal(0, sqrt( $\sigma^2$ ))
    x[1] ~ Normal( $\mu$ , sqrt( $\sigma^2$ ))
    x[2] ~ Normal( $\mu$ , sqrt( $\sigma^2$ ))
     $\sigma^2$ ,  $\mu$ 
end

model = gdemo([1.5, 2])
alg1 = PG(50, 1000)
chn1 = sample(model, alg)
alg2 = HMC(1000, 0.2, 3)
chn2 = sample(model, alg)
```

Code 2: Gaussian model with conjugate priors (inference steps breakdown)

- Non-deterministic?
 - Distributions
- Language?
 - @model macro
 - ~ notation (macro)
- Machine learning?
 - *Bayesian inference* by sampling: SMC, PG, HMC, NUTS, Gibbs ...

WHY JULIA?

- Rich collections of statistical libraries
 - StatsFuns.jl, Distributions.jl, Mamba.jl
- Meta-programming
 - Turing's compiler, i.e. @model and @~
- Coroutines
 - Particle Gibbs implementation
- Automatic differentiation (AD)
 - ForwardDiff.jl provides an easy-to-use Dual type
 - Hamiltonian Monte Carlo implementation
 - Generic typing help AD work with distributions

KEY FEATURES

- Universal probabilistic programming with an intuitive modelling interface embedded in Julia
- Support of models with discrete variables and stochastic control flows by particle filtering
- HMC sampler for differentiable distributions
- 🌀 Compositional MCMC interface
- 🌀 Resumption of MCMC chains
- ☐ More novel samplers, other inference methods, sampling from user-defined models, ...

🌀 = new releases, ☐ = next steps

$$\mu_k \sim \text{Normal}(0, 25), \sigma_k^2 \sim \text{Inv-Gamma}(2, 3), k = 1 \dots K$$

$$z_i \sim \text{Cat}(1/K), x_i \sim \text{Normal}(\mu_{z_i}, \sigma_{z_i}^2), i = 1 \dots N$$

```
using Turing
```

```
@model MoG(x, K) = begin
```

```
    μ = Vector{Real}(K)
```

```
    σ² = Vector{Real}(K)
```

```
    for i = 1:K
```

```
        μ[i] ~ Normal(0, 25)
```

```
        σ²[i] ~ InverseGamma(2, 3)
```

```
    end
```

```
    N = length(x); z = zeros{Int, N}
```

```
    for i = 1:N
```

```
        z[i] ~ Categorical(1/K*ones(K))
```

```
        x[i] ~ Normal(μ[z[i]], sqrt(σ²[z[i]]))
```

```
    end
```

```
    z, σ², μ
```

```
end
```

```
modelf = MoG([1.1, 1.0, 0.9, 2.1, 2.2], 2)
```

```
alg = Gibbs(1000, PG(50, 1, :z), HMC(1, 0.2, 3, :μ, :σ²))
```

```
chn = sample(modelf, alg)
```

Code 3: Compositional MCMC interface


```
using Turing

@model somemodel(...) = begin
    ...
end

modelf = somemodel(...)
chn1 = sample(modelf, HMC(1000, 1, 0.2, 3; save_state=true))
chn2 = sample(modelf, NUTS(1000, 0.65; resume_from=chn1))
```

Code 3: Resumption of MCMC chains

HOW TURING.JL WORKS?

Key ML techniques

- Bayesian inference
 - General framework for probabilistic modelling
- Sampling
 - Particle filtering
 - Markov chain Monte Carlo

Key Julia techniques

- Coroutines
 - Particle based samplers
- Automatic differentiation
 - Gradient based samplers

Key system components

- Model *defined by* users
 - Normal Julia program with modelling operations
- Sampler *specified by* users
 - Need to interact with model
- `VarInfo`
 - Key data structure
 - Enable interactions between models and samplers
 - Users don't see it
- Samples *returned to* users
 - Embedded in `Chain` type

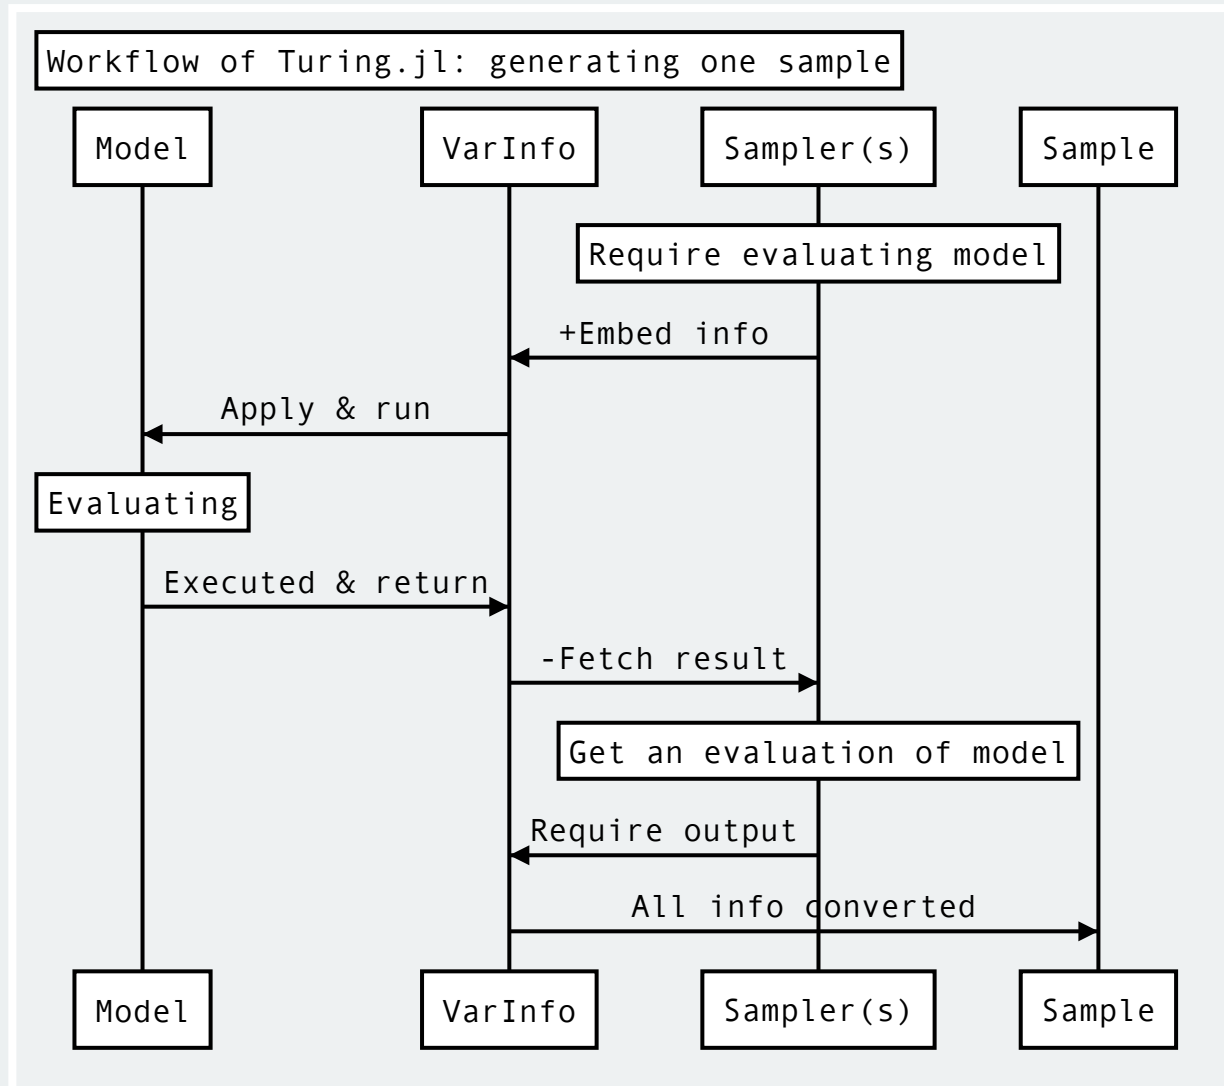


Figure 1: Workflow of Turing.jl: generating one sample

HOW PARTICLE FILTERING WORKS IN TURING.JL

- IS, SMC and PG represent each particle running in parallel using a coroutine
- Each particle is a "live" copy of the model, i.e. an ongoing execution of the model
 - Samplers need to duplicate or kill particles
 - SMC and PG do model evaluation sequentially for each observation
 - Resuming and pausing a particle is necessary
- Lead to state-of-the-art SMC and PG performance

HOW HMC WORKS IN TURING.JL

1. HMC algorithm needs the unnormalized posterior of the model and the gradient of variables
2. Execution of the model program with variables set as `ForwardDiff.Dual` gives both
 - Unnormalized posterior = real part of log-joint
 - Gradient of variables = dual parts of log-joint
3. Produce a sample candidate by simulating Hamiltonian dynamics with the leapfrog algorithm
4. Accept or reject the sample candidate

LIVE DEMO

Time to see a live Turing.jl program ...

Q&A

Any question?

THANK YOU FOR
LISTENING 😊